

Cryptography

Lecture 13

Arkady Yerukhimovich

October 9, 2024

- 1 Lecture 12 Review
- 2 Hash Functions (Chapters 5.1, 5.2)
- 3 Other Applications of Hash Functions (Chapters 5.3, 5.6)

Lecture 12 Review

- Review of MAC domain extension
- Authenticated encryption

- 1 Lecture 12 Review
- 2 Hash Functions (Chapters 5.1, 5.2)
- 3 Other Applications of Hash Functions (Chapters 5.3, 5.6)

Domain Extension for MAC (Try 4)

Starting Point

- Let $m = m_1 || m_2 || \dots || m_\ell$, where each m_i is n bits
- Let $\Pi' = (\text{Gen}', \text{Mac}', \text{Verify}')$ be an n -bit MAC

Include random message identifier in each block:

- Parse m as $m_1 || m_2 || \dots || m_{4\ell}$ with each m_i of length $n/4$
- $r \leftarrow \{0, 1\}^{n/4}$ - message id
- Compute $t_i = \text{Mac}'_k(r || 4\ell || i || m_i)$

The Problem:

This requires

- $|t| = 4\ell n$ bits
- 4ℓ calls to PRF

Question: Can we do domain extension more efficiently?

Another Way to Authenticate Long Messages

What if we could take a *digest* of a long message?

$$\begin{array}{c} m = m_1 || m_2 || \quad \dots \quad || m_\ell \\ \downarrow \\ H(m) \end{array}$$

and, then compute $t = \text{Mac}_k(H(m))$

Question

What properties would we need from H for this to be a secure Mac?

Hash Functions

A hash function is a function $H : \{0, 1\}^* \rightarrow \{0, 1\}^\ell$

Hash Functions

A hash function is a function $H : \{0, 1\}^* \rightarrow \{0, 1\}^\ell$

- *compresses* long inputs into short (fixed-length) *digests*

Hash Functions

A hash function is a function $H : \{0, 1\}^* \rightarrow \{0, 1\}^\ell$

- *compresses* long inputs into short (fixed-length) *digests*

Security:

Hash Functions

A hash function is a function $H : \{0, 1\}^* \rightarrow \{0, 1\}^\ell$

- *compresses* long inputs into short (fixed-length) *digests*

Security:

- *Second pre-image resistance*: Given m and $H(m)$, can't find m' with same hash

Hash Functions

A hash function is a function $H : \{0, 1\}^* \rightarrow \{0, 1\}^\ell$

- *compresses* long inputs into short (fixed-length) *digests*

Security:

- *Second pre-image resistance*: Given m and $H(m)$, can't find m' with same hash
- *Collision resistance*: Hard to find m, m' s.t. $H(m) = H(m')$.

A More Formal Definition

(Keyed) Hash function with output length $\ell(n)$:

A More Formal Definition

(Keyed) Hash function with output length $\ell(n)$:

- $\text{Gen}(1^n)$: Outputs a key s
- $H(s, x)$: $H^s(x)$ takes input $x \in \{0, 1\}^*$ and outputs $H^s(x) \in \{0, 1\}^{\ell(n)}$

A More Formal Definition

(Keyed) Hash function with output length $\ell(n)$:

- $\text{Gen}(1^n)$: Outputs a key s
- $H(s, x)$: $H^s(x)$ takes input $x \in \{0, 1\}^*$ and outputs $H^s(x) \in \{0, 1\}^{\ell(n)}$

Security Definition

Let $\Pi = (\text{Gen}, H)$ be a (keyed) hash function with output length ℓ .
Consider the following game between an adversary \mathcal{A} and a challenger:

Hash – $\text{Coll}_{\mathcal{A}, \Pi}(n)$

Security Definition

Let $\Pi = (\text{Gen}, H)$ be a (keyed) hash function with output length ℓ .
Consider the following game between an adversary \mathcal{A} and a challenger:

Hash – $\text{Coll}_{\mathcal{A}, \Pi}(n)$

- The challenger chooses $s \leftarrow \text{Gen}(1^n)$ and sends it to \mathcal{A}

Security Definition

Let $\Pi = (\text{Gen}, H)$ be a (keyed) hash function with output length ℓ .
Consider the following game between an adversary \mathcal{A} and a challenger:

Hash – $\text{Coll}_{\mathcal{A}, \Pi}(n)$

- The challenger chooses $s \leftarrow \text{Gen}(1^n)$ and sends it to \mathcal{A}
- \mathcal{A} outputs two strings (x, x')

Security Definition

Let $\Pi = (\text{Gen}, H)$ be a (keyed) hash function with output length ℓ .
Consider the following game between an adversary \mathcal{A} and a challenger:

Hash – $\text{Coll}_{\mathcal{A}, \Pi}(n)$

- The challenger chooses $s \leftarrow \text{Gen}(1^n)$ and sends it to \mathcal{A}
- \mathcal{A} outputs two strings (x, x')
- We say that $\text{Hash} - \text{Coll}_{\mathcal{A}, \Pi}(n) = 1$ (i.e., \mathcal{A} wins) if $H^s(x) = H^s(x')$.

Security Definition

Let $\Pi = (\text{Gen}, H)$ be a (keyed) hash function with output length ℓ .
Consider the following game between an adversary \mathcal{A} and a challenger:

Hash – Coll $_{\mathcal{A},\Pi}(n)$

- The challenger chooses $s \leftarrow \text{Gen}(1^n)$ and sends it to \mathcal{A}
- \mathcal{A} outputs two strings (x, x')
- We say that Hash – Coll $_{\mathcal{A},\Pi}(n) = 1$ (i.e., \mathcal{A} wins) if $H^s(x) = H^s(x')$.

Definition: A hash function $\Pi = (\text{Gen}, H)$ is *collision resistant* if for all PPT \mathcal{A} it holds that

$$\Pr[\text{Hash} - \text{Coll}_{\mathcal{A},\Pi}(n) = 1] \leq \text{negl}(n)$$

Security Definition

Let $\Pi = (\text{Gen}, H)$ be a (keyed) hash function with output length ℓ .
Consider the following game between an adversary \mathcal{A} and a challenger:

Hash – Coll $_{\mathcal{A},\Pi}(n)$

- The challenger chooses $s \leftarrow \text{Gen}(1^n)$ and sends it to \mathcal{A}
- \mathcal{A} outputs two strings (x, x')
- We say that Hash – Coll $_{\mathcal{A},\Pi}(n) = 1$ (i.e., \mathcal{A} wins) if $H^s(x) = H^s(x')$.

Comparison to a PRF:

Security Definition

Let $\Pi = (\text{Gen}, H)$ be a (keyed) hash function with output length ℓ .
Consider the following game between an adversary \mathcal{A} and a challenger:

Hash – Coll $_{\mathcal{A}, \Pi}(n)$

- The challenger chooses $s \leftarrow \text{Gen}(1^n)$ and sends it to \mathcal{A}
- \mathcal{A} outputs two strings (x, x')
- We say that Hash – Coll $_{\mathcal{A}, \Pi}(n) = 1$ (i.e., \mathcal{A} wins) if $H^s(x) = H^s(x')$.

Comparison to a PRF:

- The key s is given to \mathcal{A} (in PRF key had to be secret)

Security Definition

Let $\Pi = (\text{Gen}, H)$ be a (keyed) hash function with output length ℓ .
Consider the following game between an adversary \mathcal{A} and a challenger:

Hash – $\text{Coll}_{\mathcal{A}, \Pi}(n)$

- The challenger chooses $s \leftarrow \text{Gen}(1^n)$ and sends it to \mathcal{A}
- \mathcal{A} outputs two strings (x, x')
- We say that $\text{Hash} - \text{Coll}_{\mathcal{A}, \Pi}(n) = 1$ (i.e., \mathcal{A} wins) if $H^s(x) = H^s(x')$.

Comparison to a PRF:

- The key s is given to \mathcal{A} (in PRF key had to be secret)
- No oracle is necessary, \mathcal{A} has s , so can compute $H^s(\cdot)$

Comparison to a MAC:

Security Definition

Let $\Pi = (\text{Gen}, H)$ be a (keyed) hash function with output length ℓ .
Consider the following game between an adversary \mathcal{A} and a challenger:

Hash – $\text{Coll}_{\mathcal{A}, \Pi}(n)$

- The challenger chooses $s \leftarrow \text{Gen}(1^n)$ and sends it to \mathcal{A}
- \mathcal{A} outputs two strings (x, x')
- We say that $\text{Hash} - \text{Coll}_{\mathcal{A}, \Pi}(n) = 1$ (i.e., \mathcal{A} wins) if $H^s(x) = H^s(x')$.

Comparison to a PRF:

- The key s is given to \mathcal{A} (in PRF key had to be secret)
- No oracle is necessary, \mathcal{A} has s , so can compute $H^s(\cdot)$

Comparison to a MAC:

- Given $y = H^s(m)$, hard to find m' that hashes to y

Security Definition

Let $\Pi = (\text{Gen}, H)$ be a (keyed) hash function with output length ℓ .
Consider the following game between an adversary \mathcal{A} and a challenger:

Hash – $\text{Coll}_{\mathcal{A}, \Pi}(n)$

- The challenger chooses $s \leftarrow \text{Gen}(1^n)$ and sends it to \mathcal{A}
- \mathcal{A} outputs two strings (x, x')
- We say that $\text{Hash} - \text{Coll}_{\mathcal{A}, \Pi}(n) = 1$ (i.e., \mathcal{A} wins) if $H^s(x) = H^s(x')$.

Comparison to a PRF:

- The key s is given to \mathcal{A} (in PRF key had to be secret)
- No oracle is necessary, \mathcal{A} has s , so can compute $H^s(\cdot)$

Comparison to a MAC:

- Given $y = H^s(m)$, hard to find m' that hashes to y
- But, since s is public, any party can produce $(m', y' = H^s(m'))$

Hash Functions in Practice

Hash functions used in practice are a bit different:

Hash Functions in Practice

Hash functions used in practice are a bit different:

- Output length is fixed (e.g., 256-bits), not a function of n

Hash Functions in Practice

Hash functions used in practice are a bit different:

- Output length is fixed (e.g., 256-bits), not a function of n
- Usually used unkeyed

Hash Functions in Practice

Hash functions used in practice are a bit different:

- Output length is fixed (e.g., 256-bits), not a function of n
- Usually used unkeyed
- Still generally believed to be collision resistant

Hash Functions in Practice

Hash functions used in practice are a bit different:

- Output length is fixed (e.g., 256-bits), not a function of n
- Usually used unkeyed
- Still generally believed to be collision resistant

For now, we will stick to the asymptotic definition

How Secure Can a Hash Function Be?

A hash function maps arbitrary length strings to $\ell(n)$ length strings:

How Secure Can a Hash Function Be?

A hash function maps arbitrary length strings to $\ell(n)$ length strings:

- There are more possible inputs than outputs, collisions must occur

How Secure Can a Hash Function Be?

A hash function maps arbitrary length strings to $\ell(n)$ length strings:

- There are more possible inputs than outputs, collisions must occur
- If hash cuts input in half: $H : \{0, 1\}^{2^\ell} \rightarrow \{0, 1\}^\ell$, then 2^{2^ℓ} values mapped to 2^ℓ boxes, so there are very many collisions

How Secure Can a Hash Function Be?

A hash function maps arbitrary length strings to $\ell(n)$ length strings:

- There are more possible inputs than outputs, collisions must occur
- If hash cuts input in half: $H : \{0, 1\}^{2^\ell} \rightarrow \{0, 1\}^\ell$, then 2^{2^ℓ} values mapped to 2^ℓ boxes, so there are very many collisions
- How many values do I need to try before I find such a collision?

How Secure Can a Hash Function Be?

A hash function maps arbitrary length strings to $\ell(n)$ length strings:

- There are more possible inputs than outputs, collisions must occur
- If hash cuts input in half: $H : \{0, 1\}^{2\ell} \rightarrow \{0, 1\}^{\ell}$, then $2^{2\ell}$ values mapped to 2^{ℓ} boxes, so there are very many collisions
- How many values do I need to try before I find such a collision?

Birthday Paradox

How Secure Can a Hash Function Be?

A hash function maps arbitrary length strings to $\ell(n)$ length strings:

- There are more possible inputs than outputs, collisions must occur
- If hash cuts input in half: $H : \{0, 1\}^{2^\ell} \rightarrow \{0, 1\}^\ell$, then 2^{2^ℓ} values mapped to 2^ℓ boxes, so there are very many collisions
- How many values do I need to try before I find such a collision?

Birthday Paradox

How many people do you need before you expect two to have same birthday?

How Secure Can a Hash Function Be?

A hash function maps arbitrary length strings to $\ell(n)$ length strings:

- There are more possible inputs than outputs, collisions must occur
- If hash cuts input in half: $H : \{0, 1\}^{2\ell} \rightarrow \{0, 1\}^{\ell}$, then $2^{2\ell}$ values mapped to 2^{ℓ} boxes, so there are very many collisions
- How many values do I need to try before I find such a collision?

Birthday Paradox

How many people do you need before you expect two to have same birthday?

- Tempting to say 365, but this is wrong.

How Secure Can a Hash Function Be?

A hash function maps arbitrary length strings to $\ell(n)$ length strings:

- There are more possible inputs than outputs, collisions must occur
- If hash cuts input in half: $H : \{0, 1\}^{2\ell} \rightarrow \{0, 1\}^{\ell}$, then $2^{2\ell}$ values mapped to 2^{ℓ} boxes, so there are very many collisions
- How many values do I need to try before I find such a collision?

Birthday Paradox

How many people do you need before you expect two to have same birthday?

- Tempting to say 365, but this is wrong.
- Consider all pairs of people (p_i, p_j) for $i \neq j$

How Secure Can a Hash Function Be?

A hash function maps arbitrary length strings to $\ell(n)$ length strings:

- There are more possible inputs than outputs, collisions must occur
- If hash cuts input in half: $H : \{0, 1\}^{2\ell} \rightarrow \{0, 1\}^{\ell}$, then $2^{2\ell}$ values mapped to 2^{ℓ} boxes, so there are very many collisions
- How many values do I need to try before I find such a collision?

Birthday Paradox

How many people do you need before you expect two to have same birthday?

- Tempting to say 365, but this is wrong.
- Consider all pairs of people (p_i, p_j) for $i \neq j$
- There are $\binom{n}{2} = \frac{n(n-1)}{2} = O(n^2)$ such pairs

How Secure Can a Hash Function Be?

A hash function maps arbitrary length strings to $\ell(n)$ length strings:

- There are more possible inputs than outputs, collisions must occur
- If hash cuts input in half: $H : \{0, 1\}^{2\ell} \rightarrow \{0, 1\}^{\ell}$, then $2^{2\ell}$ values mapped to 2^{ℓ} boxes, so there are very many collisions
- How many values do I need to try before I find such a collision?

Birthday Paradox

How many people do you need before you expect two to have same birthday?

- Tempting to say 365, but this is wrong.
- Consider all pairs of people (p_i, p_j) for $i \neq j$
- There are $\binom{n}{2} = \frac{n(n-1)}{2} = O(n^2)$ such pairs
- For a pair, $\Pr[bday_i = bday_j] = 1/365$

How Secure Can a Hash Function Be?

A hash function maps arbitrary length strings to $\ell(n)$ length strings:

- There are more possible inputs than outputs, collisions must occur
- If hash cuts input in half: $H : \{0, 1\}^{2\ell} \rightarrow \{0, 1\}^{\ell}$, then $2^{2\ell}$ values mapped to 2^{ℓ} boxes, so there are very many collisions
- How many values do I need to try before I find such a collision?

Birthday Paradox

How many people do you need before you expect two to have same birthday?

- Tempting to say 365, but this is wrong.
- Consider all pairs of people (p_i, p_j) for $i \neq j$
- There are $\binom{n}{2} = \frac{n(n-1)}{2} = O(n^2)$ such pairs
- For a pair, $\Pr[bday_i = bday_j] = 1/365$
- After 365 pairs (28 people), expect a collision

How Secure Can a Hash Function Be?

A hash function maps arbitrary length strings to $\ell(n)$ length strings:

- There are more possible inputs than outputs, collisions must occur
- If hash cuts input in half: $H : \{0, 1\}^{2\ell} \rightarrow \{0, 1\}^{\ell}$, then $2^{2\ell}$ values mapped to 2^{ℓ} boxes, so there are very many collisions
- How many values do I need to try before I find such a collision?

Birthday Paradox

How many people do you need before you expect two to have same birthday?

- Tempting to say 365, but this is wrong.
- Consider all pairs of people (p_i, p_j) for $i \neq j$
- There are $\binom{n}{2} = \frac{n(n-1)}{2} = O(n^2)$ such pairs
- For a pair, $\Pr[bday_i = bday_j] = 1/365$
- After 365 pairs (28 people), expect a collision
- Generally, $O(2^{\ell/2})$ for output length ℓ – need ℓ large enough

Building a Hash Function

How to build a hash function:

Building a Hash Function

How to build a hash function:


- Start with a *compression function* $h^s : \{0, 1\}^{\ell'} \rightarrow \{0, 1\}^{\ell}$
 - Defined only for a fixed $\ell' > \ell$
 - We will see examples later in the course

Building a Hash Function


How to build a hash function:

- Start with a *compression function* $h^s : \{0, 1\}^{\ell'} \rightarrow \{0, 1\}^{\ell}$
 - Defined only for a fixed $\ell' > \ell$
 - We will see examples later in the course
- Extend domain from ℓ' -bit strings to arbitrary bit strings
 - This is what we will do now

Merkle-Damgård Domain Extension




Merkle-Damgård Domain Extension




- Let $h^s : \{0, 1\}^{2\ell} \rightarrow \{0, 1\}^\ell$ be a compression function

Merkle-Damgård Domain Extension




- Let $h^s : \{0, 1\}^{2\ell} \rightarrow \{0, 1\}^\ell$ be a compression function
- Given input $x \in \{0, 1\}^L$

Merkle-Damgård Domain Extension




- Let $h^s : \{0, 1\}^{2\ell} \rightarrow \{0, 1\}^\ell$ be a compression function
- Given input $x \in \{0, 1\}^L$
- Break x into ℓ -bit blocks

Merkle-Damgård Domain Extension




- Let $h^s : \{0, 1\}^{2\ell} \rightarrow \{0, 1\}^\ell$ be a compression function
- Given input $x \in \{0, 1\}^L$
- Break x into ℓ -bit blocks
- Add length L as last block

Merkle-Damgård Domain Extension




- Let $h^s : \{0, 1\}^{2\ell} \rightarrow \{0, 1\}^\ell$ be a compression function
- Given input $x \in \{0, 1\}^L$
- Break x into ℓ -bit blocks
- Add length L as last block
- Compute $H^s(x)$ as in the figure above

Merkle-Damgård Domain Extension




Proof of Collision Resistance:

Merkle-Damgård Domain Extension



Proof of Collision Resistance: Show collision in H^s gives collision in h^s .


Merkle-Damgård Domain Extension



Proof of Collision Resistance: Show collision in H^s gives collision in h^s .

Suppose $x = (x_1, \dots, x_B)$ and $x' = (x'_1, \dots, x'_B)$ collide ($H^s(x) = H^s(x')$)

Merkle-Damgård Domain Extension




Proof of Collision Resistance: Show collision in H^s gives collision in h^s .

Suppose $x = (x_1, \dots, x_B)$ and $x' = (x'_1, \dots, x'_B)$ collide ($H^s(x) = H^s(x')$)

- Case 1: $L \neq L'$

Merkle-Damgård Domain Extension




Proof of Collision Resistance: Show collision in H^s gives collision in h^s .

Suppose $x = (x_1, \dots, x_B)$ and $x' = (x'_1, \dots, x'_B)$ collide ($H^s(x) = H^s(x')$)

- Case 1: $L \neq L'$

- $h^s(z_B || L) = h^s(z'_B || L')$, but $L \neq L'$ – collision

Merkle-Damgård Domain Extension




Proof of Collision Resistance: Show collision in H^s gives collision in h^s .

Suppose $x = (x_1, \dots, x_B)$ and $x' = (x'_1, \dots, x'_B)$ collide ($H^s(x) = H^s(x')$)

- Case 1: $L \neq L'$
 - $h^s(z_B || L) = h^s(z'_B || L')$, but $L \neq L'$ – collision
- Case 2: $L = L'$

Merkle-Damgård Domain Extension




Proof of Collision Resistance: Show collision in H^s gives collision in h^s .

Suppose $x = (x_1, \dots, x_B)$ and $x' = (x'_1, \dots, x'_B)$ collide ($H^s(x) = H^s(x')$)

- Case 1: $L \neq L'$
 - $h^s(z_B || L) = h^s(z'_B || L')$, but $L \neq L'$ – collision
- Case 2: $L = L'$
 - Find largest index where inputs to h^s are different

Merkle-Damgård Domain Extension



Proof of Collision Resistance: Show collision in H^s gives collision in h^s .

Suppose $x = (x_1, \dots, x_B)$ and $x' = (x'_1, \dots, x'_{B'})$ collide ($H^s(x) = H^s(x')$)


- Case 1: $L \neq L'$

- $h^s(z_B || L) = h^s(z'_{B'} || L')$, but $L \neq L'$ – collision

- Case 2: $L = L'$

- Find largest index where inputs to h^s are different
 - Such index must exist since $x \neq x'$

Merkle-Damgård Domain Extension



Proof of Collision Resistance: Show collision in H^s gives collision in h^s .

Suppose $x = (x_1, \dots, x_B)$ and $x' = (x'_1, \dots, x'_B)$ collide ($H^s(x) = H^s(x')$)

- Case 1: $L \neq L'$
 - $h^s(z_B || L) = h^s(z'_B || L')$, but $L \neq L'$ – collision
- Case 2: $L = L'$
 - Find largest index where inputs to h^s are different
 - Such index must exist since $x \neq x'$
 - At this index, you have two different inputs to h^s that produce same output – collision

Domain Extension for MACs (Hash-and-MAC)

Building blocks

- $\Pi = (\text{Gen}, \text{Mac}, \text{Verify})$: Secure MAC for $\ell(n)$ -bit messages
- $\Pi_H = (\text{Gen}_H, H)$: CRHF with output length $\ell(n)$

Domain Extension for MACs (Hash-and-MAC)

Building blocks

- $\Pi = (\text{Gen}, \text{Mac}, \text{Verify})$: Secure MAC for $\ell(n)$ -bit messages
- $\Pi_H = (\text{Gen}_H, H)$: CRHF with output length $\ell(n)$

Construct $\Pi' = (\text{Gen}', \text{Mac}', \text{Verify}')$:

Domain Extension for MACs (Hash-and-MAC)

Building blocks

- $\Pi = (\text{Gen}, \text{Mac}, \text{Verify})$: Secure MAC for $\ell(n)$ -bit messages
- $\Pi_H = (\text{Gen}_H, H)$: CRHF with output length $\ell(n)$

Construct $\Pi' = (\text{Gen}', \text{Mac}', \text{Verify}')$:

- $\text{Gen}'(1^n)$: Choose $k \leftarrow \text{Gen}(1^n)$, $s \leftarrow \text{Gen}_H(1^n)$. Output $k' = (k, s)$

Domain Extension for MACs (Hash-and-MAC)

Building blocks

- $\Pi = (\text{Gen}, \text{Mac}, \text{Verify})$: Secure MAC for $\ell(n)$ -bit messages
- $\Pi_H = (\text{Gen}_H, H)$: CRHF with output length $\ell(n)$

Construct $\Pi' = (\text{Gen}', \text{Mac}', \text{Verify}')$:

- $\text{Gen}'(1^n)$: Choose $k \leftarrow \text{Gen}(1^n)$, $s \leftarrow \text{Gen}_H(1^n)$. Output $k' = (k, s)$
- $\text{Mac}'_{k'}(m)$: Recall that $k' = (k, s)$, Output $t = \text{Mac}_k(H^s(m))$

Domain Extension for MACs (Hash-and-MAC)

Building blocks

- $\Pi = (\text{Gen}, \text{Mac}, \text{Verify})$: Secure MAC for $\ell(n)$ -bit messages
- $\Pi_H = (\text{Gen}_H, H)$: CRHF with output length $\ell(n)$

Construct $\Pi' = (\text{Gen}', \text{Mac}', \text{Verify}')$:

- $\text{Gen}'(1^n)$: Choose $k \leftarrow \text{Gen}(1^n)$, $s \leftarrow \text{Gen}_H(1^n)$. Output $k' = (k, s)$
- $\text{Mac}'_{k'}(m)$: Recall that $k' = (k, s)$, Output $t = \text{Mac}_k(H^s(m))$
- $\text{Verify}'_{k'}(m, t)$: Output 1 if $t = \text{Mac}_k(H^s(m))$

Hash-and-MAC

Construct $\Pi' = (\text{Gen}', \text{Mac}', \text{Verify}')$:

- $\text{Gen}'(1^n)$: Choose $k \leftarrow \text{Gen}(1^n)$, $s \leftarrow \text{Gen}_H(1^n)$. Output $k' = (k, s)$
- $\text{Mac}'_{k'}(m)$: Recall that $k' = (k, s)$, Output $t = \text{Mac}_k(H^s(m))$
- $\text{Verify}'_{k'}(m, t)$: Output 1 if $t = \text{Mac}_k(H^s(m))$

Hash-and-MAC

Construct $\Pi' = (\text{Gen}', \text{Mac}', \text{Verify}')$:

- $\text{Gen}'(1^n)$: Choose $k \leftarrow \text{Gen}(1^n)$, $s \leftarrow \text{Gen}_H(1^n)$. Output $k' = (k, s)$
- $\text{Mac}'_{k'}(m)$: Recall that $k' = (k, s)$, Output $t = \text{Mac}_k(H^s(m))$
- $\text{Verify}'_{k'}(m, t)$: Output 1 if $t = \text{Mac}_k(H^s(m))$

Proof Sketch: Suppose \mathcal{A} forges a valid t on $m^* \notin Q$

Hash-and-MAC

Construct $\Pi' = (\text{Gen}', \text{Mac}', \text{Verify}')$:

- $\text{Gen}'(1^n)$: Choose $k \leftarrow \text{Gen}(1^n)$, $s \leftarrow \text{Gen}_H(1^n)$. Output $k' = (k, s)$
- $\text{Mac}'_{k'}(m)$: Recall that $k' = (k, s)$, Output $t = \text{Mac}_k(H^s(m))$
- $\text{Verify}'_{k'}(m, t)$: Output 1 if $t = \text{Mac}_k(H^s(m))$

Proof Sketch: Suppose \mathcal{A} forges a valid t on $m^* \notin Q$

- Case 1: $\exists m \in Q$ s.t. $H^s(m^*) = H^s(m)$

Hash-and-MAC

Construct $\Pi' = (\text{Gen}', \text{Mac}', \text{Verify}')$:

- $\text{Gen}'(1^n)$: Choose $k \leftarrow \text{Gen}(1^n)$, $s \leftarrow \text{Gen}_H(1^n)$. Output $k' = (k, s)$
- $\text{Mac}'_{k'}(m)$: Recall that $k' = (k, s)$, Output $t = \text{Mac}_k(H^s(m))$
- $\text{Verify}'_{k'}(m, t)$: Output 1 if $t = \text{Mac}_k(H^s(m))$

Proof Sketch: Suppose \mathcal{A} forges a valid t on $m^* \notin Q$

- Case 1: $\exists m \in Q$ s.t. $H^s(m^*) = H^s(m)$
 - Then (m, m^*) is a collision in H

Hash-and-MAC

Construct $\Pi' = (\text{Gen}', \text{Mac}', \text{Verify}')$:

- $\text{Gen}'(1^n)$: Choose $k \leftarrow \text{Gen}(1^n)$, $s \leftarrow \text{Gen}_H(1^n)$. Output $k' = (k, s)$
- $\text{Mac}'_{k'}(m)$: Recall that $k' = (k, s)$, Output $t = \text{Mac}_k(H^s(m))$
- $\text{Verify}'_{k'}(m, t)$: Output 1 if $t = \text{Mac}_k(H^s(m))$

Proof Sketch: Suppose \mathcal{A} forges a valid t on $m^* \notin Q$

- Case 1: $\exists m \in Q$ s.t. $H^s(m^*) = H^s(m)$
 - Then (m, m^*) is a collision in H
- Case 2: $\forall m \in Q, H^s(m^*) \neq H^s(m)$

Hash-and-MAC

Construct $\Pi' = (\text{Gen}', \text{Mac}', \text{Verify}')$:

- $\text{Gen}'(1^n)$: Choose $k \leftarrow \text{Gen}(1^n)$, $s \leftarrow \text{Gen}_H(1^n)$. Output $k' = (k, s)$
- $\text{Mac}'_{k'}(m)$: Recall that $k' = (k, s)$, Output $t = \text{Mac}_k(H^s(m))$
- $\text{Verify}'_{k'}(m, t)$: Output 1 if $t = \text{Mac}_k(H^s(m))$

Proof Sketch: Suppose \mathcal{A} forges a valid t on $m^* \notin Q$

- Case 1: $\exists m \in Q$ s.t. $H^s(m^*) = H^s(m)$
 - Then (m, m^*) is a collision in H
- Case 2: $\forall m \in Q$, $H^s(m^*) \neq H^s(m)$
 - Let $H^s(Q) = \{H^s(m) \mid m \in Q\}$

Hash-and-MAC

Construct $\Pi' = (\text{Gen}', \text{Mac}', \text{Verify}')$:

- $\text{Gen}'(1^n)$: Choose $k \leftarrow \text{Gen}(1^n)$, $s \leftarrow \text{Gen}_H(1^n)$. Output $k' = (k, s)$
- $\text{Mac}'_{k'}(m)$: Recall that $k' = (k, s)$, Output $t = \text{Mac}_k(H^s(m))$
- $\text{Verify}'_{k'}(m, t)$: Output 1 if $t = \text{Mac}_k(H^s(m))$

Proof Sketch: Suppose \mathcal{A} forges a valid t on $m^* \notin Q$

- Case 1: $\exists m \in Q$ s.t. $H^s(m^*) = H^s(m)$
 - Then (m, m^*) is a collision in H
- Case 2: $\forall m \in Q$, $H^s(m^*) \neq H^s(m)$
 - Let $H^s(Q) = \{H^s(m) \mid m \in Q\}$
 - Then, $H^s(m^*) \notin H^s(Q)$

Hash-and-MAC

Construct $\Pi' = (\text{Gen}', \text{Mac}', \text{Verify}')$:

- $\text{Gen}'(1^n)$: Choose $k \leftarrow \text{Gen}(1^n)$, $s \leftarrow \text{Gen}_H(1^n)$. Output $k' = (k, s)$
- $\text{Mac}'_{k'}(m)$: Recall that $k' = (k, s)$, Output $t = \text{Mac}_k(H^s(m))$
- $\text{Verify}'_{k'}(m, t)$: Output 1 if $t = \text{Mac}_k(H^s(m))$

Proof Sketch: Suppose \mathcal{A} forges a valid t on $m^* \notin Q$

- Case 1: $\exists m \in Q$ s.t. $H^s(m^*) = H^s(m)$
 - Then (m, m^*) is a collision in H
- Case 2: $\forall m \in Q, H^s(m^*) \neq H^s(m)$
 - Let $H^s(Q) = \{H^s(m) \mid m \in Q\}$
 - Then, $H^s(m^*) \notin H^s(Q)$
 - But, then \mathcal{A} has forged valid tag on new message $H^s(m^*)$

- 1 Lecture 12 Review
- 2 Hash Functions (Chapters 5.1, 5.2)
- 3 Other Applications of Hash Functions (Chapters 5.3, 5.6)

What Are Hash Functions Good For?

Properties of Hash Functions

- Produce short digest of long strings (e.g., files, etc.)
- Collision resistant – hard to find two strings that have same hash
- $H(x)$ uniquely identifies an item x

What Are Hash Functions Good For?

Properties of Hash Functions

- Produce short digest of long strings (e.g., files, etc.)
- Collision resistant – hard to find two strings that have same hash
- $H(x)$ uniquely identifies an item x

Applications of unique identifiers:

What Are Hash Functions Good For?

Properties of Hash Functions

- Produce short digest of long strings (e.g., files, etc.)
- Collision resistant – hard to find two strings that have same hash
- $H(x)$ uniquely identifies an item x

Applications of unique identifiers:

- Fingerprinting – store $H(x)$ where x is a virus

What Are Hash Functions Good For?

Properties of Hash Functions

- Produce short digest of long strings (e.g., files, etc.)
- Collision resistant – hard to find two strings that have same hash
- $H(x)$ uniquely identifies an item x

Applications of unique identifiers:

- Fingerprinting – store $H(x)$ where x is a virus
- Deduplication – check if two files (e.g., in cloud storage) are the same

What Are Hash Functions Good For?

Properties of Hash Functions

- Produce short digest of long strings (e.g., files, etc.)
- Collision resistant – hard to find two strings that have same hash
- $H(x)$ uniquely identifies an item x

Applications of unique identifiers:

- Fingerprinting – store $H(x)$ where x is a virus
- Deduplication – check if two files (e.g., in cloud storage) are the same
- Peer-to-peer file sharing – Use $H(\text{file})$ as unique identifier

Password-Based Authentication

How to use a password:

- 1 User creates password *pwd* when registering for site

Password-Based Authentication

How to use a password:

- 1 User creates password *pwd* when registering for site
- 2 Site stores (ID, pwd) on server

Password-Based Authentication

How to use a password:

- ① User creates password pwd when registering for site
- ② Site stores (ID, pwd) on server
- ③ When logging in, user enters pwd' , server compares it to pwd

Password-Based Authentication

How to use a password:

- 1 User creates password pwd when registering for site
- 2 Site stores (ID, pwd) on server
- 3 When logging in, user enters pwd' , server compares it to pwd

We have a problem

- Server stores all users' passwords
- If this file is stolen, \mathcal{A} can impersonate any user

Password-Based Authentication

How to use a password:

- 1 User creates password pwd when registering for site
- 2 Site stores (ID, pwd) on server
- 3 When logging in, user enters pwd' , server compares it to pwd

We have a problem

- Server stores all users' passwords
- If this file is stolen, \mathcal{A} can impersonate any user

Can we protect passwords even if password file is stolen?

Hash The Passwords

How to use a password:

- 1 User creates password *pwd* when registering for site

Hash The Passwords

How to use a password:

- ① User creates password pwd when registering for site
- ② Site computes $H(pwd)$, stores $(ID, H(pwd))$ on server

Hash The Passwords

How to use a password:

- ① User creates password pwd when registering for site
- ② Site computes $H(pwd)$, stores $(ID, H(pwd))$ on server
- ③ When logging in, user enters pwd' , server computes $H(pwd')$ and checks password file

Hash The Passwords

How to use a password:

- 1 User creates password pwd when registering for site
- 2 Site computes $H(pwd)$, stores $(ID, H(pwd))$ on server
- 3 When logging in, user enters pwd' , server computes $H(pwd')$ and checks password file

Pros:

- Passwords not stored in clear
- \mathcal{A} who steals password file only gets hashes of passwords

Hash The Passwords

How to use a password:

- 1 User creates password pwd when registering for site
- 2 Site computes $H(pwd)$, stores $(ID, H(pwd))$ on server
- 3 When logging in, user enters pwd' , server computes $H(pwd')$ and checks password file

Pros:

- Passwords not stored in clear
- \mathcal{A} who steals password file only gets hashes of passwords

Cons:

- \mathcal{A} can build big list of hashed passwords
- When \mathcal{A} gets files of hashed passwords, can see if any on the list

Hash The Passwords

How to use a password:

- 1 User creates password pwd when registering for site
- 2 Site computes $H(pwd)$, stores $(ID, H(pwd))$ on server
- 3 When logging in, user enters pwd' , server computes $H(pwd')$ and checks password file

Pros:

- Passwords not stored in clear
- \mathcal{A} who steals password file only gets hashes of passwords

Cons:

- \mathcal{A} can build big list of hashed passwords
- When \mathcal{A} gets files of hashed passwords, can see if any on the list
- Recovering password of any user is good enough

Hash The Passwords

How to use a password:

- 1 User creates password pwd when registering for site
- 2 Site computes $H(pwd)$, stores $(ID, H(pwd))$ on server
- 3 When logging in, user enters pwd' , server computes $H(pwd')$ and checks password file

Pros:

- Passwords not stored in clear
- \mathcal{A} who steals password file only gets hashes of passwords

Cons:

- \mathcal{A} can build big list of hashed passwords
- When \mathcal{A} gets files of hashed passwords, can see if any on the list
- Recovering password of any user is good enough
- Hashing is very fast – Billions of hashes / second

Add a Salt

How to use a password:

- 1 User creates password *pwd* when registering for site

Add a Salt

How to use a password:

- 1 User creates password pwd when registering for site
- 2 Site chooses salt $s \leftarrow \{0, 1\}^n$, computes $H(s||pwd)$ and stores $(ID, s, H(s||pwd))$

Add a Salt

How to use a password:

- 1 User creates password pwd when registering for site
- 2 Site chooses salt $s \leftarrow \{0, 1\}^n$, computes $H(s||pwd)$ and stores $(ID, s, H(s||pwd))$
- 3 When logging in, user enters pwd' , server finds s , computes $H(s||pwd')$ and checks password file

Add a Salt

How to use a password:

- 1 User creates password pwd when registering for site
- 2 Site chooses salt $s \leftarrow \{0, 1\}^n$, computes $H(s||pwd)$ and stores $(ID, s, H(s||pwd))$
- 3 When logging in, user enters pwd' , server finds s , computes $H(s||pwd')$ and checks password file

Pros:

- s is long and random, can't be predicted by \mathcal{A}

Add a Salt

How to use a password:

- 1 User creates password pwd when registering for site
- 2 Site chooses salt $s \leftarrow \{0, 1\}^n$, computes $H(s||pwd)$ and stores $(ID, s, H(s||pwd))$
- 3 When logging in, user enters pwd' , server finds s , computes $H(s||pwd')$ and checks password file

Pros:

- s is long and random, can't be predicted by \mathcal{A}
- If \mathcal{A} builds list of hashed passwords, it'll be useless if s is wrong

Add a Salt

How to use a password:

- 1 User creates password pwd when registering for site
- 2 Site chooses salt $s \leftarrow \{0, 1\}^n$, computes $H(s||pwd)$ and stores $(ID, s, H(s||pwd))$
- 3 When logging in, user enters pwd' , server finds s , computes $H(s||pwd')$ and checks password file

Pros:

- s is long and random, can't be predicted by \mathcal{A}
- If \mathcal{A} builds list of hashed passwords, it'll be useless if s is wrong
- s is unique per user, so even if he constructs list after seeing s , he can only attack one user

Add a Salt

How to use a password:

- 1 User creates password pwd when registering for site
- 2 Site chooses salt $s \leftarrow \{0, 1\}^n$, computes $H(s||pwd)$ and stores $(ID, s, H(s||pwd))$
- 3 When logging in, user enters pwd' , server finds s , computes $H(s||pwd')$ and checks password file

Pros:

- s is long and random, can't be predicted by \mathcal{A}
- If \mathcal{A} builds list of hashed passwords, it'll be useless if s is wrong
- s is unique per user, so even if he constructs list after seeing s , he can only attack one user

Takeaway

Always use a salt

Outsourced Storage

The Goal

- Client wants to store files x_1, \dots, x_n on server
- When he later retrieves file, wants to be sure it wasn't modified

Outsourced Storage

The Goal

- Client wants to store files x_1, \dots, x_n on server
- When he later retrieves file, wants to be sure it wasn't modified

Try 1: Client stores $h_1 = H(x_1), h_2 = H(x_2), \dots, h_n = H(x_n)$

Outsourced Storage

The Goal

- Client wants to store files x_1, \dots, x_n on server
- When he later retrieves file, wants to be sure it wasn't modified

Try 1: Client stores $h_1 = H(x_1), h_2 = H(x_2), \dots, h_n = H(x_n)$

- Pro: Can quickly verify each file, by recomputing $H(x_i)$

The Goal

- Client wants to store files x_1, \dots, x_n on server
- When he later retrieves file, wants to be sure it wasn't modified

Try 1: Client stores $h_1 = H(x_1), h_2 = H(x_2), \dots, h_n = H(x_n)$

- Pro: Can quickly verify each file, by recomputing $H(x_i)$
- Con: Client has to store n items

Outsourced Storage

The Goal

- Client wants to store files x_1, \dots, x_n on server
- When he later retrieves file, wants to be sure it wasn't modified

Try 1: Client stores $h_1 = H(x_1), h_2 = H(x_2), \dots, h_n = H(x_n)$

- Pro: Can quickly verify each file, by recomputing $H(x_i)$
- Con: Client has to store n items

Try 2: Client stores $h = H(x_1 || x_2 || \dots || x_n)$

Outsourced Storage

The Goal

- Client wants to store files x_1, \dots, x_n on server
- When he later retrieves file, wants to be sure it wasn't modified

Try 1: Client stores $h_1 = H(x_1), h_2 = H(x_2), \dots, h_n = H(x_n)$

- Pro: Can quickly verify each file, by recomputing $H(x_i)$
- Con: Client has to store n items

Try 2: Client stores $h = H(x_1 || x_2 || \dots || x_n)$

- Pro: Only ℓ bits of storage needed on client

The Goal

- Client wants to store files x_1, \dots, x_n on server
- When he later retrieves file, wants to be sure it wasn't modified

Try 1: Client stores $h_1 = H(x_1), h_2 = H(x_2), \dots, h_n = H(x_n)$

- Pro: Can quickly verify each file, by recomputing $H(x_i)$
- Con: Client has to store n items

Try 2: Client stores $h = H(x_1 || x_2 || \dots || x_n)$

- Pro: Only ℓ bits of storage needed on client
- Con: Must download all files to verify even a single one

Outsourced Storage

The Goal

- Client wants to store files x_1, \dots, x_n on server
- When he later retrieves file, wants to be sure it wasn't modified

Try 1: Client stores $h_1 = H(x_1), h_2 = H(x_2), \dots, h_n = H(x_n)$

- Pro: Can quickly verify each file, by recomputing $H(x_i)$
- Con: Client has to store n items


Try 2: Client stores $h = H(x_1 || x_2 || \dots || x_n)$

- Pro: Only ℓ bits of storage needed on client
- Con: Must download all files to verify even a single one


Question: Can we get solution that achieves both?

- Low storage on the client
- Low communication to verify a file is correct

Merkle Tree




Building a Merkle Tree




- C has files x_{000}, \dots, x_{111} , computes $H_{000} = H(x_{000}), \dots$

Building a Merkle Tree




- C has files x_{000}, \dots, x_{111} , computes $H_{000} = H(x_{000}), \dots$
- C computes $H_{00} = H(H_{000}, H_{001}), H_{01} = H(H_{010}, H_{011}), \dots$

Building a Merkle Tree




- C has files x_{000}, \dots, x_{111} , computes $H_{000} = H(x_{000}), \dots$
- C computes $H_{00} = H(H_{000}, H_{001}), H_{01} = H(H_{010}, H_{011}), \dots$
- C does the same for all levels: $H_0 = H(H_{00}, H_{01}), H_R = H(H_0, H_1)$

Building a Merkle Tree




- C has files x_{000}, \dots, x_{111} , computes $H_{000} = H(x_{000}), \dots$
- C computes $H_{00} = H(H_{000}, H_{001}), H_{01} = H(H_{010}, H_{011}), \dots$
- C does the same for all levels: $H_0 = H(H_{00}, H_{01}), H_R = H(H_0, H_1)$
- C stores value H_R at the root and uploads all files to S

Verifying File x_{001} :




- C has H_R , downloads x_{001} and wants to check it's correct

Verifying File x_{001} :




- C has H_R , downloads x_{001} and wants to check it's correct
- S sends C all sibling hashes along the path to root: H_{000}, H_{01}, H_1

Verifying File x_{001} :




- C has H_R , downloads x_{001} and wants to check it's correct
- S sends C all sibling hashes along the path to root: H_{000}, H_{01}, H_1
- C computes H_{001} and can now recompute up to H_R (has both inputs for every node in path)

Verifying File x_{001} :




- C has H_R , downloads x_{001} and wants to check it's correct
- S sends C all sibling hashes along the path to root: H_{000}, H_{01}, H_1
- C computes H_{001} and can now recompute up to H_R (has both inputs for every node in path)
- C checks that computed value is equal to H_R , accepts if so

Merkle Tree Facts




- Allows verified outsourced storage

Merkle Tree Facts



- Allows verified outsourced storage
- C only needs to store a single hash value, H_R

Merkle Tree Facts



- Allows verified outsourced storage
- C only needs to store a single hash value, H_R
- Proof consists of $O(\log n)$ hash values